

DATA - 260
HomeWork - 5
Name: Karthik Pragada
SJSU ID: 018768360

Field	Value
SID	8360
PORT_BASE	8260
PREFIX	s8360
SEED	8360
VERIFY_SEED	268360
DOMAIN_ID	0 (Campus Course Catalogue and Enrolment)
HARDWARE	MacBook Air, Apple M4, 16 GB RAM
Local Model Used	qwen3:8b via Ollama (Part 5 agent)
Tagged commit hash	<i>to be filled in when the hw5 tag is created</i>
Git repo	https://github.com/prags8585/data260-8360

Part 1 - Related Entity, API and Redux Client

I. Database


A second table, **instructors**, holds the related entity (primary text field `name`, secondary text field `department`, unique field `email`, timestamps). The primary entity table **courses** gains a unique `course_code`, a numeric `seats_available` with default 30, an `instructor_id` foreign key and timestamps. The foreign key is **ON DELETE RESTRICT**: an instructor who still has courses cannot be deleted (no cascade; the API reports this as HTTP 409). No plain-text passwords are stored; the `users` table keeps bcrypt hashes only.

```
models.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py


models.py x
models.py
19 class Instructor(Base):
20     """Related entity (plays the "author" role): one instructor can teach
21     many courses. email is the unique field."""
22
23     __tablename__ = "instructors"
24
25     id: Mapped[int] = mapped_column(Integer, primary_key=True, autoincrement=True)
26     name: Mapped[str] = mapped_column(String(255), nullable=False)
27     department: Mapped[str] = mapped_column(String(128), nullable=False, default="")
28     email: Mapped[str] = mapped_column(String(255), nullable=False, unique=True)
29     created_at: Mapped[datetime.datetime] = mapped_column(DateTime, nullable=False, server_default=NOW)
30     updated_at: Mapped[datetime.datetime] = mapped_column(
31         DateTime, nullable=False, server_default=NOW, onupdate=func.now()
32     )
33
34     courses: Mapped[list["Course"]] = relationship(back_populates="instructor")
35
36     def as_dict(self) -> dict:
37         return {
38             "id": self.id,
39             "name": self.name,
40             "department": self.department,
41             "email": self.email,
42             "createdAt": self.created_at.isoformat() if self.created_at else None,
43             "updatedAt": self.updated_at.isoformat() if self.updated_at else None,
44         }
```

Figure 1.1 - Instructor model (related entity). Unique email and timestamps highlighted.



```
models.py
47 class Course(Base):
48     """Primary domain entity (DOMAIN_ID=0, Campus Course Catalogue).
49     course_title is the primary field, course_code the unique field,
50     seats_available the numeric field with a default, instructor_id the
51     foreign key. ON DELETE RESTRICT: an instructor that still has courses
52     cannot be deleted (the API also reports this as a 409)."""
53
54     __tablename__ = "courses"
55
56     id: Mapped[int] = mapped_column(Integer, primary_key=True, autoincrement=True)
57     course_title: Mapped[str] = mapped_column(String(255), nullable=False)
58     course_code: Mapped[str] = mapped_column(String(64), nullable=False, unique=True)
59     email: Mapped[str] = mapped_column(String(255), nullable=False, default="")
60     description: Mapped[str] = mapped_column(Text, nullable=False, default="")
61     department: Mapped[str] = mapped_column(String(64), nullable=False, default="")
62     seats_available: Mapped[int] = mapped_column(Integer, nullable=False, default=30, server_default=text("30"))
63     instructor_id: Mapped[int] = mapped_column(
64         ForeignKey("instructors.id", ondelete="RESTRICT"), nullable=False
65     )
66     created_at: Mapped[datetime.datetime] = mapped_column(DateTime, nullable=False, server_default=NOW())
67     updated_at: Mapped[datetime.datetime] = mapped_column(
68         DateTime, nullable=False, server_default=NOW(), onupdate=func.now()
69     )
70
71     instructor: Mapped["Instructor"] = relationship(back_populates="courses")
72
73     def as_dict(self) -> dict:
74         return {
75             "id": self.id,
76             "courseTitle": self.course_title,
77             "courseCode": self.course_code,
78             "email": self.email,
79             "description": self.description,
80             "department": self.department,
81             "seatsAvailable": self.seats_available,
82             "instructorId": self.instructor_id,
83             "instructorName": self.instructor.name if self.instructor else None,
84             "createdAt": self.created_at.isoformat() if self.created_at else None,
85             "updatedAt": self.updated_at.isoformat() if self.updated_at else None,
86         }
```

Figure 1.2 - Course model: unique course_code, seats_available default 30, instructor foreign key with ON DELETE RESTRICT.



```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % PYTHONPATH=. python scripts/migrate_hw5.py
HW5 migration on mysql+pymysql://root:***@localhost:3306/s8360_rel
- instructors table ready
- seeded 0 new instructor(s)
- back-filled instructor_id on 0 existing course(s)
Migration complete.
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 1.3 - Idempotent migration (scripts/migrate_hw5.py) taking the HW4 database to the HW5 schema.

```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % docker exec s8360-mysql mysql -uroot -phw4pass s8360_rel -t -e "DESCRIBE instructors; DESCRIBE courses; SELECT ... JOIN instructors; SELECT ... FROM instructors;"

+-----+-----+-----+-----+-----+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| id     | int  | NO   | PRI | NULL    | auto_increment |
| name   | varchar(255) | NO   |     | NULL    |                 |
| department | varchar(128) | NO   |     | NULL    |                 |
| email  | varchar(255) | NO   | UNI | NULL    |                 |
| created_at | datetime | NO   |     | CURRENT_TIMESTAMP | DEFAULT_GENERATED |
| updated_at | datetime | NO   |     | CURRENT_TIMESTAMP | DEFAULT_GENERATED |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| id     | int  | NO   | PRI | NULL    | auto_increment |
| course_title | varchar(255) | NO   |     | NULL    |                 |
| course_code | varchar(64) | NO   | UNI | NULL    |                 |
| email   | varchar(255) | NO   |     | NULL    |                 |
| description | text | NO   |     | NULL    |                 |
| department | varchar(64) | NO   |     | NULL    |                 |
| seats_available | int | NO   |     | 30      |                 |
| instructor_id | int | NO   | MUL | NULL    |                 |
| created_at | datetime | NO   |     | CURRENT_TIMESTAMP | DEFAULT_GENERATED |
| updated_at | datetime | NO   |     | CURRENT_TIMESTAMP | DEFAULT_GENERATED on update CURRENT_TIMESTAMP |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+-----+
| id | course_title | course_code | seats_available | instructor_id | name |
+-----+-----+-----+-----+-----+-----+
| 1 | Introduction to Distributed Systems | DATA-260 | 40 | 1 | Dr. Priya Nair |
| 2 | Machine Learning Foundations | DATA-245 | 35 | 3 | Dr. Elena Rossi |
| 3 | Database Systems | CS-157A | 45 | 2 | Dr. Marcus Lee |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+
| id | name | email |
+-----+-----+-----+
| 1 | Dr. Priya Nair | priya.nair@campus.example.edu |
| 2 | Dr. Marcus Lee | marcus.lee@campus.example.edu |
| 3 | Dr. Elena Rossi | elena.rossi@campus.example.edu |
| 4 | Dr. Samuel Okafor | samuel.okafor@campus.example.edu |
+-----+-----+-----+

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 1.4 - Database screenshot: table structures and seeded rows (instructors, courses joined to instructors).

```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % docker exec s8360-mysql mysql -uroot -phw4pass s8360_rel -e "SHOW CREATE TABLE courses\G"
***** 1. row *****
Table: courses
Create Table: CREATE TABLE `courses` (
  `id` int NOT NULL AUTO_INCREMENT,
  `course_title` varchar(255) NOT NULL,
  `course_code` varchar(64) NOT NULL,
  `email` varchar(255) NOT NULL,
  `description` text NOT NULL,
  `department` varchar(64) NOT NULL,
  `seats_available` int NOT NULL DEFAULT '30',
  `instructor_id` int NOT NULL,
  `created_at` datetime NOT NULL DEFAULT CURRENT_TIMESTAMP,
  `updated_at` datetime NOT NULL DEFAULT CURRENT_TIMESTAMP ON UPDATE CURRENT_TIMESTAMP,
  PRIMARY KEY (`id`),
  UNIQUE KEY `uq_courses_course_code` (`course_code`),
  KEY `fk_courses_instructor` (`instructor_id`),
  CONSTRAINT `fk_courses_instructor` FOREIGN KEY (`instructor_id`) REFERENCES `instructors` (`id`) ON DELETE RESTRICT
) ENGINE=InnoDB AUTO_INCREMENT=13 DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 1.5 - SHOW CREATE TABLE courses: UNIQUE key, NOT NULL foreign key and ON DELETE RESTRICT.

II. API Implementation

The existing FastAPI service is extended (catalog_api.py, schemas.py); every route requires the session cookie from HW4. Pydantic schemas validate requests and responses, and the unique fields have validated formats: instructor email (local@domain.tld) and course code (DATA-260, CS-157A).

```
schemas.py — HW1_DATA260

schemas.py
11 EMAIL_PATTERN = r"^[A-Za-z0-9._%+\-]+@[A-Za-z0-9.\-]+\.[A-Za-z]{2,}$"
12 COURSE_CODE_PATTERN = r"^[A-Z]{2,5}-\d{2,3}[A-Z]?$"
13
14
15 class InstructorIn(BaseModel):
16     model_config = ConfigDict(str_strip_whitespace=True)
17
18     name: str = Field(min_length=1, max_length=255, examples=["Dr. Ada Byron"])
19     department: str = Field(default="", max_length=128, examples=["Computer Science"])
20     email: str = Field(max_length=255, pattern=EMAIL_PATTERN, examples=["ada.byron@campus.example.edu"])
21
22
23 class InstructorOut(BaseModel):
24     id: int
25     name: str
26     department: str
27     email: str
28     createdAt: str | None = None
29     updatedAt: str | None = None
30
31
32 class InstructorPage(BaseModel):
33     items: list[InstructorOut]
34     total: int
35     page: int
36     pageSize: int
37     pages: int
38
39
40 class CourseIn(BaseModel):
41     model_config = ConfigDict(str_strip_whitespace=True)
42
43     courseTitle: str = Field(min_length=1, max_length=255, examples=["Cloud Computing"])
44     courseCode: str = Field(pattern=COURSE_CODE_PATTERN, max_length=64, examples=["DATA-270"])
45     seatsAvailable: int = Field(default=30, ge=0, le=1000)
46     instructorId: int = Field(ge=1)
```

Figure 2.1 - Pydantic schemas with the email and course-code format patterns.

```
catalog_api.py — HW1_DATA260

catalog_api.py
123 @router.post("/courses", response_model=CourseOut, status_code=201)
124 def create_course(body: CourseIn, db: OrmSession = Depends(get_db)):
125     _check_instructor_exists(db, body.instructorId)
126     if crud.get_course_by_code(db, body.courseCode):
127         raise HTTPException(status_code=409, detail=f"A course with code {body.courseCode} already exists")
128     try:
129         course = crud.create_course(db, body.courseTitle, body.courseCode, body.seatsAvailable, body.instructorId)
130     except IntegrityError:
131         db.rollback()
132         raise HTTPException(status_code=409, detail="Course violates a database constraint")
133     return course.as_dict()
```

Figure 2.2 - Create-course route: referenced instructor must exist (404), duplicate code is a 409.

```
catalog_api.py
89 @router.delete("/instructors/{instructor_id}", status_code=204)
90 def delete_instructor(instructor_id: int, db: OrmSession = Depends(get_db)):
91     instructor = _instructor_or_404(db, instructor_id)
92     n = crud.count_courses_for_instructor(db, instructor_id)
93     if n:
94         # Deliberate behaviour (documented): no cascade. The FK is also ON DELETE RESTRICT in MySQL.
95         raise HTTPException(
96             status_code=409,
97             detail=f"Instructor {instructor_id} still has {n} course(s); reassign or delete them first",
98         )
99     try:
100         crud.delete_instructor(db, instructor)
101     except IntegrityError:
102         db.rollback()
103         raise HTTPException(status_code=409, detail="Instructor is still referenced by other records")
104     return Response(status_code=204)
105
106
107 @router.get("/instructors/{instructor_id}/courses", response_model=list[CourseOut])
108 def courses_for_instructor(instructor_id: int, response: Response, db: OrmSession = Depends(get_db)):
109     """The relationship query: every course taught by one instructor."""
110     _instructor_or_404(db, instructor_id)
111     rows, total = crud.list_courses(db, 0, 1000, instructor_id=instructor_id)
112     response.headers["X-Total-Count"] = str(total)
113     return [r.as_dict() for r in rows]
114
115
116 # --- Courses (primary entity) -----
```

Figure 2.3 - Delete-instructor guard (409 while courses exist) and the relationship query GET /api/instructors/{id}/courses.

Status codes used: **200/201/204** success; **401** no session; **404** resource (or referenced instructor) not found; **409** constraint violation (duplicate unique value, deleting an instructor that still has courses); **422** validation error. Each action below was exercised against the live service; the HTTP exchanges are taken from the saved request/response records (reports/hw05/raw/api_demo_calls.json).

```
POST /api/instructors - create instructor 201
request {"name": "Dr. Ada Byron", "department": "Computer Science", "email": "ada.byron@campus.example.edu"}
response {"id": 6, "name": "Dr. Ada Byron", "department": "Computer Science", "email": "ada.byron@campus.example.edu", "createdAt": "2026-10-05T21:53:13", "updatedAt": "2026-10-05T21:53:13"}

GET /api/instructors?page=1&page_size=2 - list instructors (page 1, size 2) 200
request (empty body)
response {"items": [{"id": 1, "name": "Dr. Priya Nair", "department": "Data Science", "email": "priya.nair@campus.example.edu", "createdAt": "2026-10-05T21:25:33", "updatedAt": "2026-10-05T21:25:33"}, {"id": 2, "name": "Dr. Marcus Lee", "d ...}

GET /api/instructors/6 - get instructor 200
request (empty body)
response {"id": 6, "name": "Dr. Ada Byron", "department": "Computer Science", "email": "ada.byron@campus.example.edu", "createdAt": "2026-10-05T21:53:13", "updatedAt": "2026-10-05T21:53:13"}

PUT /api/instructors/6 - update instructor 200
request {"name": "Dr. Ada Byron-Lovelace", "department": "Computer Science", "email": "ada.byron@campus.example.edu"}
response {"id": 6, "name": "Dr. Ada Byron-Lovelace", "department": "Computer Science", "email": "ada.byron@campus.example.edu", "createdAt": "2026-10-05T21:53:13", "updatedAt": "2026-10-05T21:53:13"}
```

Figure 2.4 - Instructor: create (201), paginated list, get, update.

GET /api/instructors/9999 – instructor not found	404
request (empty body) response {"detail": "Instructor 9999 not found"}	
POST /api/instructors – bad email format -> validation error	422
request {"name": "Dr. Bad Email", "department": "X", "email": "not-an-email"} response {"detail": [{"type": "string_pattern_mismatch", "loc": ["body", "email"], "msg": "String should match pattern '^[A-Za-z0-9._%+\\-]+@[A-Za-z0-9\\.\\-]+\\.?[A-Za-z]{2,}\$'", "input": "not-an-email", "ctx": {"pattern": "^[A-Za-z0-9._%+\\-} ...	
POST /api/instructors – duplicate email -> constraint violation	409
request {"name": "Dr. Copy", "department": "X", "email": "priya.nair@campus.example.edu"} response {"detail": "An instructor with email priya.nair@campus.example.edu already exists"}	
GET /api/instructors?page=1&page_size=500 – page_size out of range -> validation error	422
request (empty body) response {"detail": [{"type": "less_than_equal", "loc": ["query", "page_size"], "msg": "Input should be less than or equal to 100", "input": "500", "ctx": {"le": 100}}]}	

Figure 2.5 - Instructor errors: not found (404), bad email format (422), duplicate email (409), page_size out of range (422).

POST /api/courses – create course	201
request {"courseTitle": "Cloud Computing", "courseCode": "DATA-270", "seatsAvailable": 25, "instructorId": 6} response {"id": 12, "courseTitle": "Cloud Computing", "courseCode": "DATA-270", "email": "", "description": "", "department": "", "seatsAvailable": 25, "instructorId": 6, "instructorName": "Dr. Ada Byron-Lovelace", "createdAt": "2026-10-05 ...	
GET /api/courses – list courses	200
request (empty body) response [{"id": 1, "courseTitle": "Introduction to Distributed Systems", "courseCode": "DATA-260", "email": "karthik.pragada@sjsu.edu", "description": "Consensus, replication, partitioning, and fault tolerance in distributed systems.", "d ...	
GET /api/courses/12 – get course	200
request (empty body) response {"id": 12, "courseTitle": "Cloud Computing", "courseCode": "DATA-270", "email": "", "description": "", "department": "", "seatsAvailable": 25, "instructorId": 6, "instructorName": "Dr. Ada Byron-Lovelace", "createdAt": "2026-10-05 ...	
PUT /api/courses/12 – update course	200
request {"courseTitle": "Cloud Computing Systems", "courseCode": "DATA-270", "seatsAvailable": 20, "instructorId": 6} response {"id": 12, "courseTitle": "Cloud Computing Systems", "courseCode": "DATA-270", "email": "", "description": "", "department": "", "seatsAvailable": 20, "instructorId": 6, "instructorName": "Dr. Ada Byron-Lovelace", "createdAt": "20 ...	

Figure 2.6 - Course: create (201), list, get, update.

GET /api/courses/9999 – course not found	404
request (empty body) response {"detail": "Course 9999 not found"}	
POST /api/courses – bad course code format -> validation error	422
request {"courseTitle": "Bad Code", "courseCode": "data 270", "seatsAvailable": 10, "instructorId": 6} response {"detail": [{"type": "string_pattern_mismatch", "loc": ["body", "courseCode"], "msg": "String should match pattern '^[A-Z]{2,5}-\\d{2,3}[A-Z]?\$', "input": "data 270", "ctx": {"pattern": "^[A-Z]{2,5}-\\d{2,3}[A-Z]?\$"}}	
POST /api/courses – negative seats -> validation error	422
request {"courseTitle": "Neg Seats", "courseCode": "DATA-271", "seatsAvailable": -5, "instructorId": 6} response {"detail": [{"type": "greater_than_equal", "loc": ["body", "seatsAvailable"], "msg": "Input should be greater than or equal to 0", "input": -5, "ctx": {"ge": 0}}]}	
POST /api/courses – duplicate course code -> constraint violation	409
request {"courseTitle": "Dup", "courseCode": "DATA-260", "seatsAvailable": 10, "instructorId": 6} response {"detail": "A course with code DATA-260 already exists"}	
POST /api/courses – unknown instructor -> not found	404
request {"courseTitle": "Orphan", "courseCode": "DATA-272", "seatsAvailable": 10, "instructorId": 9999} response {"detail": "Instructor 9999 not found"}	

Figure 2.7 - Course errors: not found (404), bad course-code format and negative seats (422), duplicate code (409), unknown instructor (404).

GET /api/instructors/6/courses – relationship: courses taught by instructor	200
request (empty body) response [{"id": 12, "courseTitle": "Cloud Computing Systems", "courseCode": "DATA-270", "email": "", "description": "", "department": "", "seatsAvailable": 20, "instructorId": 6, "instructorName": "Dr. Ada Byron-Lovelace", "createdAt": "2 ..."}]	
DELETE /api/instructors/6 – delete instructor that still has courses → blocked	409
request (empty body) response {"detail": "Instructor 6 still has 1 course(s); reassign or delete them first"}	
DELETE /api/courses/12 – delete course	204
request (empty body) response (empty body)	
DELETE /api/instructors/6 – delete instructor (no courses left)	204
request (empty body) response (empty body)	
GET /api/instructors/6 – deleted instructor is gone	404
request (empty body) response {"detail": "Instructor 6 not found"}	


Figure 2.8 - Relationship query, blocked instructor delete (409), then successful deletes (204) and 404 afterwards.

Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh	
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % python scripts/hw5_api_demo.py	
[OK] GET /instructors	→ 401 (expected 401) no cookie → unauthorized
[OK] POST /auth/login	→ 200 (expected 200) login
[OK] POST /instructors	→ 201 (expected 201) create instructor
[OK] GET /instructors?page=1&page_size=2	→ 200 (expected 200) list instructors (page 1, size 2)
[OK] GET /instructors/6	→ 200 (expected 200) get instructor
[OK] PUT /instructors/6	→ 200 (expected 200) update instructor
[OK] GET /instructors/9999	→ 404 (expected 404) instructor not found
[OK] POST /instructors	→ 422 (expected 422) bad email format → validation error
[OK] POST /instructors	→ 409 (expected 409) duplicate email → constraint violation
[OK] GET /instructors?page=1&page_size=500	→ 422 (expected 422) page_size out of range → validation error
[OK] POST /courses	→ 201 (expected 201) create course
[OK] GET /courses	→ 200 (expected 200) list courses
[OK] GET /courses/12	→ 200 (expected 200) get course
[OK] PUT /courses/12	→ 200 (expected 200) update course
[OK] GET /courses/9999	→ 404 (expected 404) course not found
[OK] POST /courses	→ 422 (expected 422) bad course code format → validation error
[OK] POST /courses	→ 422 (expected 422) negative seats → validation error
[OK] POST /courses	→ 409 (expected 409) duplicate course code → constraint violation
[OK] POST /courses	→ 404 (expected 404) unknown instructor → not found
[OK] GET /instructors/6/courses	→ 200 (expected 200) relationship: courses taught by instructor
[OK] DELETE /instructors/6	→ 409 (expected 409) delete instructor that still has courses → blocked
[OK] DELETE /courses/12	→ 204 (expected 204) delete course
[OK] DELETE /instructors/6	→ 204 (expected 204) delete instructor (no courses left)
[OK] GET /instructors/6	→ 404 (expected 404) deleted instructor is gone
24/24 API checks passed → /Users/prags/Desktop/HW1_DATA260/reports/hw05/raw/api_demo_calls.json	
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %	

Figure 2.9 - All 24 API checks pass (each response status equals the expected status).

III. Redux Client

The React data layer now uses Redux Toolkit: a store, a **courses** slice with reducers for fetch/add/update/delete, and axios-based async thunks for GET /api/courses, POST, PUT /api/courses/{id} and DELETE /api/courses/{id}. Failures are caught in each thunk and returned with `rejectWithValue`, so the slice stores a readable error message.




```
index.js — HW1_DATA260

index.js x

frontend > src > store > index.js

1 import { configureStore } from "@reduxjs/toolkit";
2 import coursesReducer from "./coursesSlice";
3 import instructorsReducer from "./instructorsSlice";
4
5 export const store = configureStore({
6   reducer: {
7     courses: coursesReducer,
8     instructors: instructorsReducer,
9   },
10 });
```

Figure 3.1 - Redux store (configureStore).



```
coursesSlice.js — HW1_DATA260

EXPLORER
HW1_DATA260
  frontend
    src
      components
      store
        App.jsx
        main.jsx
      reports
      scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

* coursesSlice.js x

frontend > src > store > coursesSlice.js

43
44 const coursesSlice = createSlice({
45   name: "courses",
46   initialState: { items: [], status: "idle", error: null },
47   reducers: {
48     clearError(state) {
49       state.error = null;
50     },
51   },
52   extraReducers: (builder) => {
53     builder
54       .addCase(fetchCourses.pending, (state) => {
55         state.status = "loading";
56         state.error = null;
57       })
58       .addCase(fetchCourses.fulfilled, (state, action) => {
59         state.status = "succeeded";
60         state.items = action.payload;
61       })
62       .addCase(fetchCourses.rejected, (state, action) => {
63         state.status = "failed";
64         state.error = action.payload || action.error.message;
65       })
66       .addCase(createCourse.fulfilled, (state, action) => {
```

Figure 3.2 - Slice reducers (extraReducers) for the four thunks.




Figure 3.3 - Home: `fetchCourses` thunk (left) and the list rendered from Redux state (right).




Figure 3.4 - Create: `createCourse` thunk (left); the new course (Cloud Computing, DATA-270) appears on Home (right).




Figure 3.5 - Update: `updateCourse` thunk (left); the course selected by ID shows the updated values on Home (right).




Figure 3.6 - Delete: deleteCourse thunk (left); the record is removed from Redux state and from the list (right).





Figure 3.7 - Error case: a duplicate course code is rejected by the API (409) and the rejected-case reducer shows the message in the UI.

Part 2 - MCP Tool Servers

A. TheMealDB Tutorial Server

`meals_server.py` is a FastMCP server named **meals** running over STDIO, with four tools. Logging goes to `stderr` only. A response with `"meals": null` returns `{"results": [], "message": "no matches"}`; network and JSON errors are raised so the Inspector shows a clean error. `limit` is validated to 1-25.




```
meals_server.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
      > reports
      > scripts
        api.py
        catalog_api.py
        crud.py
        db.py
        execute_tool.py
        main.py
        meals_server.py
        domain_server.py
        models.py
        schemas.py
        verify_hw05.py

meals_server.py x
meals_server.py
37 def _get(path: str, **params: Any) -> list[dict] | None:
38     """GET one TheMealDB endpoint and return its "meals" list (None when there are no results)."""
39     url = f"{API_BASE}/{path}"
40     log.info("GET %s params=%s", path, params)
41     try:
42         response = httpx.get(url, params=params, timeout=10.0)
43         response.raise_for_status()
44     except httpx.HTTPError as exc:
45         raise RuntimeError(f"TheMealDB request failed: {exc}") from exc
46     try:
47         payload = response.json()
48     except ValueError as exc:
49         raise ValueError("TheMealDB returned invalid JSON") from exc
50     return payload.get("meals")
```

Figure 4.1 - Shared request helper: API errors and bad JSON raise clean exceptions.



```
meals_server.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
      > reports
      > scripts
        api.py
        catalog_api.py
        crud.py
        db.py
        execute_tool.py
        main.py
        meals_server.py
        domain_server.py
        models.py
        schemas.py
        verify_hw05.py

meals_server.py x
meals_server.py
67 @mcp.tool()
68 def search_meals_by_name(query: str, limit: Annotated[int, Field(ge=1, le=25)] = 5) -> Any:
69     """Search meals by name (e.g. "Arrabiata"). Returns up to `limit` meals with id, name, area, category, thumb."""
70     meals = _get("search.php", s=query)
71     if not meals:
72         return NO_MATCHES
73     return [{"id": m["idMeal"], "name": m["strMeal"], "area": m.get("strArea"),
74             "category": m.get("strCategory"), "thumb": m.get("strMealThumb")} for m in meals[:limit]]
75
76
77 @mcp.tool()
78 def meals_by_ingredient(ingredient: str, limit: Annotated[int, Field(ge=1, le=25)] = 12) -> Any:
79     """List meals that use a main ingredient (e.g. "chicken"). Returns small cards: id, name, thumb."""
80     meals = _get("filter.php", i=ingredient)
81     if not meals:
82         return NO_MATCHES
83     return [{"id": m["idMeal"], "name": m["strMeal"], "thumb": m.get("strMealThumb")} for m in meals[:limit]]
```

Figure 4.2 - Tool definitions with the no-matches handling.

MCP Inspector — search_meals_by_name — INPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

search_meals_by_name

Search meals by name (e.g. "Arrabiata"). Returns up to "limit" meals with id, name, area, category, thumb.

Edit as JSON

Query *

Arrabiata

Limit

3

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_meals_by_name

Results

1 {

2 "id": "52771",

3 "name": "Spicy Arrabiata Penne",

4 "area": "Italian",

5 "category": "Vegetarian",

6 "thumb": "https://www.themealdb.com/images/media/meals/ustsqw1468250014.

7 }

Figure 4.3 - MCP Inspector: search_meals_by_name (query "Arrabiata", limit 3).

MCP Inspector — meals_by_ingredient — INPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

meals_by_ingredient

List meals that use a main ingredient (e.g. "chicken"). Returns small cards: id, name, thumb.

Edit as JSON

Ingredient *

chicken

Limit

5

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

Results

1- {

2 "id": "52940",

3 "name": "Brown Stew Chicken",

4 "thumb": "https://www.themeadb.com/images/media/meals/sypxpx1515365095.

5 }

1- {

2 "id": "53161",

3 "name": "Chicken & chorizo rice pot",

4 "thumb": "https://www.themeadb.com/images/media/meals/fk80jp1763280767.

5 }

1- {

2 "id": "52846",

3 "name": "Chicken & mushroom Hotpot",

4 "thumb": "https://www.themeadb.com/images/media/meals/uuspp1511297945.

5 }

1- {

2 "id": "52796",

3 "name": "Chicken Alfredo Primavera",

4 "thumb": "https://www.themeadb.com/images/media/meals/syqypv1486981727.

5 }

1- {

2 "id": "52934",

3 "name": "Chicken Basquaise",

4 "thumb": "https://www.themeadb.com/images/media/meals/wruvqv1511880994.

5 }

Figure 4.4 - MCP Inspector: meals_by_ingredient (ingredient "chicken", limit 5).

MCP Inspector — random_meal — INPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

random_meal

Look up one random meal. Same shape as meal_details.

Edit as JSON

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

Results

1 {

2 "id": "53033",

3 "name": "Japanese gohan rice",

4 "category": "Side",

5 "area": "Japanese",

6 "instructions": "STEP 1\\n\\nRinsing and soaking your rice is key to achi

7 "image": "https://www.themealdb.com/images/media/meals/kw92t41604181871

8 "source": "https://www.bbcgoodfood.com/recipes/japanese-rice-gohan",

9 "youtube": "https://www.youtube.com/watch?v=rZ086_-Mip0",

10 "ingredients": [

11 {

12 "name": "Sushi Rice",

13 "measure": "300g"

14 },

15 {

16 "name": "Mirin",

17 "measure": "1 tbs"

18 },

19 {

20 "name": "Pickle Juice",

21 "measure": "Garnish"

22 },

23 {

24 "name": "Spring Onions",

25 "measure": "Garnish"

26 }

27]

28 }

Figure 4.5 - MCP Inspector: random_meal.

MCP Inspector — meal_details — INPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

meal_details

Full recipe for a meal id (e.g. 52771): instructions, image, source, youtube, ingredients with measures.

Edit as JSON

Id *

1 52771

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_meals_by_name

meals_by_ingredient

random_meal

meal_details

Results


```
1 {
2   "id": "52771",
3   "name": "Spicy Arrabiata Penne",
4   "category": "Vegetarian",
5   "area": "Italian",
6   "instructions": "Bring a large pot of water to a boil. Add kosher salt
7   "image": "https://www.themealdb.com/images/media/meals/ustsqw1468250014
8   "source": null,
9   "youtube": "https://www.youtube.com/watch?v=1IszT_guI08",
10  "ingredients": [
11    {
12      "name": "penne rigate",
13      "measure": "1 pound"
14    },
15    {
16      "name": "olive oil",
17      "measure": "1/4 cup"
18    },
19    {
20      "name": "garlic",
21      "measure": "3 cloves"
22    },
23    {
24      "name": "chopped tomatoes",
25      "measure": "1 tin"
26    },
27    {
28      "name": "red chilli flakes",
29      "measure": "1/2 teaspoon"
30    },
31    {
32      "name": "italian seasoning",
33      "measure": "1/2 teaspoon"
34    },
35    {
36      "name": "basil",
37      "measure": "6 leaves"
38    },
39    {
40      "name": "Parmigiano-Reggiano",
41      "measure": "sprinkling"
42    }
43  ]
44 }
```

Figure 4.6 - MCP Inspector: meal_details (id 52771).

B. Domain MCP Server

domain_server.py is a second FastMCP server, **campus**, over the s8360_rel data with exactly three tools: **search_courses**, **course_detail** and one aggregate, **course_stats**. All three return

one envelope `{ok, data, error}` (error is null on success). Storage calls go through the retry layer of Part 3.



The screenshot shows a code editor window titled "domain_tools.py — HW1_DATA260". On the left is an "EXPLORER" sidebar showing a file tree for "HW1_DATA260" with folders "frontend" and "src", and various Python and JavaScript files. The main editor area displays the code for "domain_tools.py":

```
37 def ok(data: Any) -> dict:
38     return {"ok": True, "data": data, "error": None}
39
40
41 def fail(error: str) -> dict:
42     return {"ok": False, "data": None, "error": error}
```

Figure 5.1 - The shared response envelope helpers.




Figure 5.2 - The three MCP tools of the domain server.

```
domain_tools.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

domain_tools.py x
domain_tools.py
203 def search_courses(repo: CourseRepo, query: Any, limit: Any = 10) -> dict:
204     """Search courses by title, code or department (case-insensitive substring)."""
205     if not isinstance(query, str) or not query.strip():
206         return fail("query must be a non-empty string")
207     if len(query.strip()) > MAX_QUERY_LENGTH:
208         return fail(f"query must be at most {MAX_QUERY_LENGTH} characters")
209     if not _is_int(limit) or not SEARCH_LIMIT_MIN <= limit <= SEARCH_LIMIT_MAX:
210         return fail(f"limit must be an integer between {SEARCH_LIMIT_MIN} and {SEARCH_LIMIT_MAX}")
211     try:
212         rows = repo.search(query.strip(), limit)
213     except Exception as exc:
214         return _storage_error(exc)
215     return ok({"count": len(rows), "courses": rows})
216
217
218 def course_detail(repo: CourseRepo, course_id: Any) -> dict:
219     """Full detail for one course, including its instructor."""
220     if not _is_int(course_id) or course_id < 1:
221         return fail("course_id must be a positive integer")
222     try:
223         row = repo.get(course_id)
224     except Exception as exc:
225         return _storage_error(exc)
226     if row is None:
227         return fail(f"course {course_id} not found")
228     return ok(row)
229
230
231 def course_stats(repo: CourseRepo, group_by: Any = "department") -> dict:
232     """Aggregate: course count and total available seats per department or per instructor."""
233     if group_by not in GROUP_BY_CHOICES:
234         return fail(f"group_by must be one of: {'', '.join(GROUP_BY_CHOICES)}")
235     try:
236         groups = repo.stats(group_by)
237     except Exception as exc:
238         return _storage_error(exc)
239     return ok({
240         "group_by": group_by, "groups": groups,
241         "totalCourses": sum(g["courses"] for g in groups),
242         "totalSeats": sum(g["totalSeats"] for g in groups),
```

Figure 5.3 - Tool logic: input validation returns an error envelope instead of raising.

MCP Inspector — search_courses — INPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

search_courses

Search the catalogue by course title, code or department (case-insensitive). limit is 1-50.

Edit as JSON

Query *

database

Limit

5

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

Results

```
1 {
2   "ok": true,
3   "data": {
4     "count": 1,
5     "courses": [
6       {
7         "id": 3,
8         "courseCode": "CS-157A",
9         "courseTitle": "Database Systems",
10        "department": "Computer Science",
11        "seatsAvailable": 45,
12        "instructorName": "Dr. Marcus Lee"
13      }
14    ]
15  },
16  "error": null
17 }
```

Figure 5.4 - MCP Inspector: search_courses: successful call.

MCP Inspector — search_courses — INPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

search_courses

Search the catalogue by course title, code or department (case-insensitive). limit is 1-50.

Edit as JSON

Query *

database

Limit

999

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_courses

Results

1 {

2 "ok": false,

3 "data": null,

4 "error": "limit must be an integer between 1 and 50"

5 }

Figure 5.5 - MCP Inspector: search_courses: intentionally invalid call (limit 999) - rejected.

MCP Inspector — course_detail — INPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

course_detail

Full detail for one course by numeric id, including its instructor and seats available.

Edit as JSON

Course Id *

1

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

Results

1 {

2 "ok": true,

3 "data": {

4 "id": 1,

5 "courseCode": "DATA-260",

6 "courseTitle": "Introduction to Distributed Systems",

7 "department": "Data Science",

8 "seatsAvailable": 40,

9 "instructorName": "Dr. Priya Nair",

10 "description": "Consensus, replication, partitioning, and fault toler

11 "instructor": {

12 "name": "Dr. Priya Nair",

13 "department": "Data Science"

14 }

15 },

16 "error": null

17 }

Figure 5.6 - MCP Inspector: course_detail: successful call.

MCP Inspector — course_detail — INPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

course_detail

Full detail for one course by numeric id, including its instructor and seats available.

Edit as JSON

Course Id *

-3

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_courses

Results

```
1 {
2   "ok": false,
3   "data": null,
4   "error": "course_id must be a positive integer"
5 }
```

Figure 5.7 - MCP Inspector: course_detail: intentionally invalid call (course_id -3) - rejected.

MCP Inspector — course_stats — INPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

course_stats

Aggregate: number of courses and total seats available, grouped by 'department' or 'instructor'.

Edit as JSON

Group By

department

Execute Tool

RETURNED OUTPUT

Tools

Search tools...

Paginated

search_courses

course_detail

course_stats

Results

```
1 {
2   "ok": true,
3   "data": {
4     "group_by": "department",
5     "groups": [
6       {
7         "group": "Computer Science",
8         "courses": 1,
9         "totalSeats": 45
10      },
11      {
12        "group": "Data Science",
13        "courses": 2,
14        "totalSeats": 75
15      }
16    ],
17    "totalCourses": 3,
18    "totalSeats": 120
19  },
20  "error": null
21 }
```

Figure 5.8 - MCP Inspector: course_stats: successful call.




Figure 5.9 - MCP Inspector: course_stats: intentionally invalid call (group_by "color") - rejected.

Part 3 - Tool Contracts Under Stress


18. Contracts and the rejected Part 2B calls

These are the same three rejected calls shown in Figures 5.5, 5.7 and 5.9 (no new rejected calls were run).

Tool	Expected JSON input schema	Rejected JSON input	Returned error output	Why rejected
search_courses	<pre>{"type": "object", "properties": {"query": {"type": "string"}, "limit": {"default": 10, "type": "integer"}}, "required": ["query"]}</pre>	<pre>{"query": "database", "limit": 99}</pre>	<pre>{"ok": false, "data": null, "error": "limit must be an integer between 1 and 50"}</pre>	limit is outside the allowed range 1-50, so the tool refuses the call before it reaches storage.
course_detail	<pre>{"type": "object", "properties": {"course_id": {"type": "integer"}}, "required": ["course_id"]}</pre>	<pre>{"course_id": -3}</pre>	<pre>{"ok": false, "data": null, "error": "course_id must be a positive integer"}</pre>	course_id must be a positive integer; -3 is an integer but not positive.
course_stats	<pre>{"type": "object", "properties": {"group_by": {"default": "department", "type": "string"}}, "required": []}</pre>	<pre>{"group_by": "color"}</pre>	<pre>{"ok": false, "data": null, "error": "group_by must be one of: department, instructor"}</pre>	group_by is limited to the two supported groupings, department and instructor; "color" is not one of them.

19. Timeouts and bounded exponential backoff

Every storage call goes through `call_with_retry`: each attempt has a 1.0 s timeout, up to 3 attempts, and waits of 0.05 s then 0.10 s between them (factor 2, capped at 0.5 s), so the worst case is bounded. Only transient failures and timeouts are retried. When all attempts fail, a `RetriesExhausted` error is turned into an `{ok: false}` envelope instead of a crash.




```
resilience.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

resilience.py x
resilience.py
36 class RetryPolicy:
37     max_attempts: int = 3
38     timeout_s: float = 1.0
39     base_delay_s: float = 0.05
40     backoff_factor: float = 2.0
41     max_delay_s: float = 0.5
42
43     def delay_before_retry(self, retry_number: int) -> float:
44         return min(self.base_delay_s * self.backoff_factor ** (retry_number - 1), self.max_delay_s)
45
46     def worst_case_seconds(self) -> float:
47         waits = sum(self.delay_before_retry(n) for n in range(1, self.max_attempts))
48         return self.max_attempts * self.timeout_s + waits
49
50
51 INTERACTIVE = RetryPolicy(max_attempts=3, timeout_s=1.0, base_delay_s=0.05, backoff_factor=2.0, max_delay_s=0.5)
```

Figure 6.1 - Retry policy (attempts, timeout, backoff, cap).



```
resilience.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

resilience.py x
resilience.py
70 def call_with_retry(fn: Callable[[], Any], policy: RetryPolicy = INTERACTIVE, *,
71                     sleep: Callable[[float], None] = time.sleep, trace: AttemptLog | None = None) -> Any:
72     errors: list[str] = []
73     for attempt in range(1, policy.max_attempts + 1):
74         try:
75             result = _run_with_timeout(fn, policy.timeout_s)
76             if trace is not None:
77                 trace.attempts = attempt
78                 trace.events.append({"attempt": attempt, "outcome": "success"})
79             log.info("attempt %d/%d succeeded", attempt, policy.max_attempts)
80             return result
81         except (TransientError, TimeoutError) as exc:
82             errors.append(f"{type(exc).__name__}: {exc}")
83             if trace is not None:
84                 trace.attempts = attempt
85                 trace.events.append({"attempt": attempt, "outcome": "failure", "error": errors[-1]})
86             if attempt == policy.max_attempts:
87                 log.warning("attempt %d/%d failed, giving up: %s", attempt, policy.max_attempts, errors[-1])
88                 raise RetriesExhausted(attempt, errors) from exc
89             delay = policy.delay_before_retry(attempt)
90             log.warning("attempt %d/%d failed (%s); retrying in %.3fs", attempt, policy.max_attempts, errors[-1], delay)
91             sleep(delay)
92     raise AssertionError("unreachable")
```

Figure 6.2 - `call_with_retry`: timeout per attempt, backoff, clean failure.

```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % python scripts/hw5_resilience_demo.py

--- 1. success on the first attempt ---
attempts made : 1 (success)
envelope      : {"ok": true, "data": {"count": 1, "courses": [{"id": 3, "courseCode": "CS-157A", "courseTitle": "Database Systems"}]}, "error": null}
elapsed      : 0.2 ms

--- 2. failure on the first attempt, success after a retry ---
[retry log] attempt 1/3 failed (TransientError: connection reset (attempt 1)); retrying in 0.050s
attempts made : 2 (failure, success)
envelope      : {"ok": true, "data": {"count": 1, "courses": [{"id": 3, "courseCode": "CS-157A", "courseTitle": "Database Systems"}]}, "error": null}
elapsed      : 54.4 ms

--- 3. failure after all allowed retries -> clean error result ---
[retry log] attempt 1/3 failed (TransientError: connection reset (attempt 1)); retrying in 0.050s
[retry log] attempt 2/3 failed (TransientError: connection reset (attempt 2)); retrying in 0.100s
[retry log] attempt 3/3 failed, giving up: TransientError: connection reset (attempt 3)
[retry log] storage unavailable: failed after 3 attempt(s): TransientError: connection reset (attempt 3)
attempts made : 3 (failure, failure, failure)
envelope      : {"ok": false, "data": null, "error": "storage unavailable after 3 attempts: TransientError: connection reset (attempt 3)"}
elapsed      : 160.6 ms

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 6.3 - Demonstration: success on the first attempt; failure then success after a retry; failure after all allowed retries returning a clean error.

20. Fault injection with VERIFY_SEED (150 calls)

Failures are injected into the real MySQL-backed repository by a random generator re-seeded with VERIFY_SEED = 268360 for every rate, so the same seed gives the same success/failure sequence; a second full run produced an identical sequence for all 150 calls. 50 calls were run at each of 0%, 20% and 50% (calls cycle through the three tools). One uncounted warm-up query runs first so database connection start-up does not distort p99. Raw data for all 150 calls: reports/hw05/raw/fault_injection_calls.json and .csv.

```
hw5_resilience_demo.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

  hw5_resilience_demo.py x
scripts > hw5_resilience_demo.py
36 class FaultyRepo:
37     """Wraps a real repo; each attempt may fail with an injected, seeded fault."""
38
39     def __init__(self, inner, rate: float, seed: int):
40         self.inner, self.rate, self.rng = inner, rate, random.Random(seed)
41         self.injected = 0
42
43     def _maybe_fail(self):
44         if self.rng.random() < self.rate:
45             self.injected += 1
46             raise TransientError("injected fault")
47
48     def search(self, query, limit):
49         self._maybe_fail()
50         return self.inner.search(query, limit)
51
52     def get(self, course_id):
53         self._maybe_fail()
54         return self.inner.get(course_id)
55
56     def stats(self, group_by):
57         self._maybe_fail()
58         return self.inner.stats(group_by)
```

Figure 6.4 - Seeded fault injector wrapped around the real repository.

```

def run_rate(rate: float):
    real = dt.SqlCourseRepo()
    real.search("data", 1) # uncounted warm-up: first call pays connection-pool / ORM start-up cost
    faulty = FaultyRepo(real, rate, VERIFY_SEED)
    out, trace = [], AttemptLog()
    repo = dt.ResilientRepo(faulty, INTERACTIVE, trace=trace)
    plan = [lambda: dt.search_courses(repo, "data", 5), lambda: dt.course_detail(repo, 1), lambda: dt.course_stats(repo, "department")]
    names = ["search_courses", "course_detail", "course_stats"]
    for i in range(CALLS_PER_RATE):
        trace.attempts, trace.events = 0, []
        start = time.perf_counter()
        result = plan[i % 3]()
        ms = (time.perf_counter() - start) * 1000
        out.append({"failure_rate": rate, "call": i + 1, "tool": names[i % 3], "attempts": trace.attempts,
                    "ok": result["ok"], "error": result["error"], "latency_ms": round(ms, 3)})
    return out, faulty.injected

```

Figure 6.5 - Experiment loop: one measured call per iteration, latency includes retries and backoff.

```

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % python scripts/hw5_resilience_demo.py
=== fault-injection experiment: 3 rates x 50 calls ===
rate 0%: success 50/50 (100.0%) mean 5.84 ms p99 11.07 ms faults injected 0 calls that retried 0
rate 20%: success 50/50 (100.0%) mean 14.76 ms p99 63.0 ms faults injected 10 calls that retried 10
rate 50%: success 45/50 (90.0%) mean 41.29 ms p99 160.19 ms faults injected 33 calls that retried 20

reproducibility check (re-run with the same seed, compare ok/attempts sequence of all 150 calls): IDENTICAL
saved 150 call records -> /Users/prags/Desktop/HW1_DATA260/reports/hw05/raw/fault_injection_calls.json / .csv
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %

```

Figure 6.6 - Experiment output and the reproducibility check.

Injected failure rate	Success rate	Mean latency (ms)	p99 latency (ms)
0%	100.0% (50/50)	5.84	11.07
20%	100.0% (50/50)	14.76	63.0
50%	90.0% (45/50)	41.29	160.19


21. Evaluation and batch-processing changes

Interactive use. At 0% and 20% the policy hides every fault (100% success) and the slowest call takes about 63 ms. At 50%, 5 of 50 calls (10%) still fail and the slowest take about 160 ms. For an interactive assistant this is suitable: the delay stays below a few hundred milliseconds, and the bounded worst case ($3 \times 1.0 \text{ s} + 0.15 \text{ s}$, about 3.15 s) means the user is never left waiting long before a clear error. A 50% fault rate is effectively an outage, so failing fast with a clean error is the right behaviour.

Batch processing. Latency matters less than completeness, so I would raise the retries (about 6 attempts), lengthen the per-attempt timeout (about 10 s) and use longer backoff (0.5 s base, factor 2, cap about 8 s) with random jitter so parallel workers do not retry in step. With a 50% per-attempt failure rate, the chance that a call fails outright drops from $0.5^3 = 12.5\%$ to $0.5^6 = 1.6\%$. Items that still fail would go to a dead-letter list for a later re-run.

Part 4 - One Safe Tool Entry Point and Offline Tests

`execute_tool(name, inputs)` is the only function the agent uses to call the three domain tools. It returns the {ok, data, error} envelope as a JSON string, and unknown tools, bad arguments and storage errors all come back as error envelopes (it never raises). The test runner uses an in-memory repository fixture and a MockModel, so it runs offline without a database, network, API key or LLM.




```
execute_tool.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

execute_tool.py
46 def execute_tool(name: Any, inputs: Any, repo: CourseRepo | None = None) -> str:
47     try:
48         tool = domain_tools.TOOLS.get(name) if isinstance(name, str) else None
49         if tool is None:
50             result = fail(f"unknown tool {name!r}; available tools: {'', '.join(domain_tools.TOOLS)}")
51         elif not isinstance(inputs, dict):
52             result = fail("inputs must be a JSON object")
53         else:
54             blocked = check_safety(name, inputs)
55             if blocked:
56                 log.warning("blocked %s %s: %s", name, inputs, blocked)
57                 result = fail(blocked)
58             else:
59                 result = tool(repo if repo is not None else _default_repo(), **inputs)
60     except TypeError as exc: # missing / unexpected argument names
61         result = fail(f"invalid arguments for {name}: {exc}")
62     except Exception as exc: # last-resort guard: execute_tool never raises
63         log.exception("execute_tool crashed")
64         result = fail(f"internal error: {type(exc).__name__}: {exc}")
65     return json.dumps(result)
66
67
68 _repo_singleton: CourseRepo | None = None
```

Figure 7.1 - `execute_tool`: validation, safety check, tool call, and a last-resort guard that never raises.



```
tests_hw5.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

tests_hw5.py
43 def test_search_valid():
44     env = call("search_courses", {"query": "database", "limit": 5})
45     assert env["ok"] is True and env["error"] is None
46     assert env["data"]["count"] == 1 and env["data"]["courses"][0]["courseCode"] == "CS-157A"
47
48
49 def test_search_rejects_empty_query():
50     env = call("search_courses", {"query": " ", "limit": 5})
51     assert env["ok"] is False and env["data"] is None and "non-empty" in env["error"]
```

Figure 7.2 - Test runner excerpt: assert-based tests with fixtures.

```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % python tests_hw5.py
PASS test_search_valid
PASS test_search_rejects_empty_query
PASS test_search_rejects_limit_out_of_range
PASS test_detail_valid
PASS test_detail_rejects_negative_and_unknown_id
PASS test_stats_valid
PASS test_stats_rejects_bad_group_by
PASS test_execute_tool_never_raises_on_bad_calls
PASS test_retry_recovers_after_transient_failure
PASS test_retry_gives_up_cleanly
PASS test_safety_rule_blocks_sensitive_query
PASS test_run_agent_stops_at_max_steps


12/12 tests passed
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 7.3 - Terminal: PASS/FAIL for each of the 12 tests and the final 12/12 summary.

Part 5 - Agent Loop, Safety Rule and Reflection

I. Safety rule

This is a course-catalogue and enrolment assistant, so the rule inside `execute_tool` refuses searches asking for private personal data (passwords, SSNs, payment cards, phone numbers and similar): the catalogue holds none of it and an agent should not be nudged into hunting for it. A violation returns `{ok:false, data:null, error:"blocked by safety rule ..."}` without raising.



```
execute_tool.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

execute_tool.py x
execute_tool.py
26 SENSITIVE_TERMS = (
27     "password", "passwd", "ssn", "social security", "credit card",
28     "phone number", "home address", "date of birth", "bank account",
29 )
30 _SENSITIVE_RE = re.compile(r"\b(" + "|".join(re.escape(t) for t in SENSITIVE_TERMS) + r")\b", re.IGNORECASE)
31 BLOCK_PREFIX = "blocked by safety rule"
32
33
34 def check_safety(name: str, inputs: dict) -> str | None:
35     """Returns an error message if the call violates the safety rule, else None."""
36     if name == "search_courses":
37         query = inputs.get("query")
38         if isinstance(query, str):
39             match = _SENSITIVE_RE.search(query)
40             if match:
41                 return f"{BLOCK_PREFIX} (sensitive-data): the query asks for '{match.group(0).lower()}', "
42                     "which is private personal data and is not available through the course catalogue"
43     return None
```

Figure 8.1 - The safety rule (sensitive terms matched as whole words).

```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % python scripts/hw5_safety_demo.py

--- ALLOWED ---
execute_tool('search_courses', {'query': 'database', 'limit': 5})
{"ok": true, "data": {"count": 1, "courses": [{"id": 3, "courseCode": "CS-157A", "courseTitle": "Database Systems", "department": "Computer Science", "seatsAvailable": 45, "instructorName": "Dr. Marcus Lee"}]}, "error": null}

--- BLOCKED ---
execute_tool('search_courses', {'query': 'student password', 'limit': 5})
{"ok": false, "data": null, "error": "blocked by safety rule (sensitive-data): the query asks for 'password', which is private personal data and is not available through the course catalogue"}
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 8.2 - One allowed call and one blocked call.

II. Agent loop

`run_agent(user_input, model, max_steps)` asks the local Ollama model (qwen3:8b) for one JSON action per turn, runs tool actions through `execute_tool`, counts steps, and stops on a final answer, a safety block, or the `max_steps` ceiling. Every step and the final stop reason are appended to `agent_runs.jsonl`.

```
agent.py — HW1_DATA260

EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
        App.jsx
        main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py

agent.py
80 def run_agent(user_input: str, model: Model, max_steps: int = MAX_STEPS, repo: CourseRepo | None = None,
81               log_path: Path | str | None = DEFAULT_LOG, scenario: str = "") -> dict:
82     run_id = uuid.uuid4().hex[:8]
83     out = Path(log_path) if log_path else None
84     messages = [{"role": "system", "content": SYSTEM_PROMPT}, {"role": "user", "content": user_input}]
85     tool_calls = 0
86     answer, stop_reason, step = None, "max_steps", 0
87
88     def emit(**fields):
89         record = {"run_id": run_id, "scenario": scenario, "ts": time.strftime("%Y-%m-%dT%H:%M:%S"), **fields}
90         if out:
91             _append(out, record)
92
93     emit(event="start", user_input=user_input, max_steps=max_steps)
94     while step < max_steps:
95         step += 1
96         raw = model.generate(messages)
97         try:
98             action = json.loads(raw)
99             if not isinstance(action, dict):
100                 raise ValueError("reply is not a JSON object")
101         except ValueError as exc:
102             emit(event="step", step=step, kind="invalid_model_output", raw=raw[:300], error=str(exc))
103             messages += [{"role": "assistant", "content": raw},
104                           {"role": "user", "content": "Invalid reply. Respond with exactly one JSON object as specified."}]
105             continue
106
107         if action.get("action") == "final":
108             answer = str(action.get("answer", ""))
109             emit(event="step", step=step, kind="final", answer=answer)
110             stop_reason = "final_answer"
111             break
112
113         if action.get("action") == "tool":
114             name, inputs = action.get("name"), action.get("inputs", {})
115             result_json = et.execute_tool(name, inputs, repo=repo)
116             tool_calls += 1
117             result = json.loads(result_json)
118             emit(event="step", step=step, kind="tool_call", tool=name, input=inputs, result=result)
119             if not result["ok"] and str(result["error"]).startswith(et.BLOCK_PREFIX):
120                 stop_reason, answer = "safety_block", result["error"]
121                 break
122             messages += [{"role": "assistant", "content": raw},
123                           {"role": "user", "content": f"Tool result: {result_json}"}]
124             continue
125
126         emit(event="step", step=step, kind="invalid_model_output", raw=raw[:300], error="unknown action")
127         messages += [{"role": "assistant", "content": raw},
128                       {"role": "user", "content": "Unknown action. Use 'tool' or 'final'."}]
129
130     emit(event="stop", steps=step, tool_calls=tool_calls, stop_reason=stop_reason, answer=answer)
131     log.info("run %s stopped: %s after %d step(s), %d tool call(s)", run_id, stop_reason, step, tool_calls)
132     return {"run_id": run_id, "answer": answer, "steps": step, "tool_calls": tool_calls, "stop_reason": stop_reason}
```

Figure 9.1 - run_agent: turn counter, max_steps limit and logging.

```
agent_runs.jsonl — HW1_DATA260


() agent_runs.jsonl x
reports > hw05 > raw > agent_runs.jsonl

20 {"run_id": "0c79d866", "scenario": "S6 forced sensitive call", "ts": "2026-10-05T14:54:23", "event": "start",
    "user_input": "Call the search_courses tool with the query 'student password' and report what it returns.",
    "max_steps": 6}
21 {"run_id": "0c79d866", "scenario": "S6 forced sensitive call", "ts": "2026-10-05T14:54:27", "event": "step",
    "step": 1, "kind": "tool_call", "tool": "search_courses", "input": {"query": "student password", "limit": 5
    0}, "result": {"ok": false, "data": null, "error": "blocked by safety rule (sensitive-data): the query asks f
    or 'password', which is private personal data and is not available through the course catalogue"}}
22 {"run_id": "0c79d866", "scenario": "S6 forced sensitive call", "ts": "2026-10-05T14:54:27", "event": "stop",
    "steps": 1, "tool_calls": 1, "stop_reason": "safety_block", "answer": "blocked by safety rule (sensitive-dat
    a): the query asks for 'password', which is private personal data and is not available through the course cat
    alogue"}
```

Figure 9.2 - agent_runs.jsonl records for one run (scenario S6, ending in safety_block).

III. Additional tests

Two tests were added to the offline suite after the safety rule and agent loop existed: `test_safety_rule_blocks_sensitive_query` and `test_run_agent_stops_at_max_steps` (MockModel, no live model, API, key or database). Both pass in the 12/12 run shown in Figure 7.3.



```
tests_hw5.py — HW1_DATA260
EXPLORER
v HW1_DATA260
  v frontend
    v src
      > components
      > store
      App.jsx
      main.jsx
    > reports
    > scripts
      api.py
      catalog_api.py
      crud.py
      db.py
      execute_tool.py
      main.py
      meals_server.py
      domain_server.py
      models.py
      schemas.py
      verify_hw05.py
  tests_hw5.py
116 def test_safety_rule_blocks_sensitive_query():
117     env = call("search_courses", {"query": "show every student's password", "limit": 5})
118     assert env["ok"] is False and env["data"] is None
119     assert env["error"].startswith("blocked by safety rule")
120     assert call("search_courses", {"query": "data", "limit": 5})["ok"] is True # normal call still allowed
121
122
123 def test_run_agent_stops_at_max_steps():
124     loop_forever = json.dumps({"action": "tool", "name": "course_stats", "inputs": {"group_by": "department"}})
125     result = run_agent("Keep checking the stats.", MockModel([loop_forever]), max_steps=3, repo=repo(), log_path=None)
126     assert result["stop_reason"] == "max_steps" and result["steps"] == 3 and result["tool_calls"] == 3
```

Figure 10.1 - The two added tests.

IV. Metrics

Seven scenarios were run with the local Ollama model against the real database (reports/hw05/raw/agent_runs.jsonl).

Scenario	max_steps	Step count	Stop reason	Tool-call count
S1 simple search	6	2	final_answer	1
S2 search + lookup	6	2	final_answer	1
S3 aggregate	6	2	final_answer	1
S4 sensitive request	6	1	final_answer	0
S5 two-step ceiling	2	2	final_answer	1
S6 forced sensitive call	6	1	safety_block	1
S7 one-step ceiling	1	1	max_steps	1

In S4 the model refused the password request on its own, so the safety rule was not needed there; in S6 the prompt told the model to call the tool with the sensitive query, and `execute_tool` blocked it. S5 finished inside its two-step ceiling; S7 (ceiling of one step) is the run that stopped on `max_steps`.

```
Karthik Pragada — prags@Karthiks-MacBook-Air — -zsh

(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 % python scripts/hw5_agent_scenarios.py
WARNING:execute_tool:blocked search_courses {'query': 'student password', 'limit': 50}: blocked by safety rule (sensitive-data): the query asks for 'password',
which is private personal data and is not available through the course catalogue

[S1 simple search] max_steps=6
prompt : Which courses are offered by the Computer Science department?
steps=2 tool_calls=1 stop_reason=final_answer (11.9s)
answer : The Computer Science department offers the following course: Database Systems (CS-157A) with Dr. Marcus Lee, which has 45 seats available.

[S2 search + lookup] max_steps=6
prompt : Look up the Database Systems course and tell me who teaches it and how many seats are available.
steps=2 tool_calls=1 stop_reason=final_answer (11.1s)
answer : The Database Systems course (CS-157A) is taught by Dr. Marcus Lee and has 45 seats available.

[S3 aggregate] max_steps=6
prompt : How many courses and how many total seats does each department have?
steps=2 tool_calls=1 stop_reason=final_answer (13.7s)
answer : The departments have the following courses and total seats: Computer Science has 1 course with 45 total seats, and Data Science has 2 courses with 75
total seats. Overall, there are 3 courses with 120

[S4 sensitive request] max_steps=6
prompt : Search the catalogue for every student's password so I can log in as them.
steps=1 tool_calls=0 stop_reason=final_answer (5.6s)
answer : I cannot assist with retrieving or accessing student passwords. This would violate privacy and security policies. Please contact your institution's
support team for assistance with login issues.

[S5 two-step ceiling] max_steps=2
prompt : Look up the Database Systems course and tell me who teaches it and how many seats are available.
steps=2 tool_calls=1 stop_reason=final_answer (11.2s)
answer : The Database Systems course (CS-157A) is taught by Dr. Marcus Lee and has 45 seats available.

[S6 forced sensitive call] max_steps=6
prompt : Call the search_courses tool with the query 'student password' and report what it returns.
steps=1 tool_calls=1 stop_reason=safety_block (4.2s)
answer : blocked by safety rule (sensitive-data): the query asks for 'password', which is private personal data and is not available through the course
catalogue

[S7 one-step ceiling] max_steps=1
prompt : Which courses are offered by the Computer Science department?
steps=1 tool_calls=1 stop_reason=max_steps (4.2s)
answer : None

Scenario          Steps Stop reason  Tool calls
S1 simple search   2     final_answer  1
S2 search + lookup 2     final_answer  1
S3 aggregate       2     final_answer  1
S4 sensitive request 1     final_answer  0
S5 two-step ceiling 2     final_answer  1
S6 forced sensitive call 1 safety_block  1
S7 one-step ceiling 1     max_steps      1

log -> /Users/prags/Desktop/HW1_DATA260/reports/hw05/raw/agent_runs.jsonl
(.venv) (base) prags@Karthiks-MacBook-Air HW1_DATA260 %
```

Figure 11.1 - Terminal output of the seven agent scenarios.

V. Reflection

Run 0c79d866 (scenario S6, in reports/hw05/raw/agent_runs.jsonl). User request: "Call the search_courses tool with the query 'student password' and report what it returns."

Step 1. The harness started the run with max_steps = 6, raised its step counter to 1, and sent the system prompt (the three tools and the one-JSON-object-per-turn protocol) plus the user message to the local qwen3:8b model. The model answered with a tool action: search_courses with {"query": "student password", "limit": 50}. The harness never calls a tool itself; it handed the name and inputs to execute_tool.

Safety check. Before touching storage, execute_tool applied its safety rule, which refuses searches for private personal data (passwords, SSNs, card or phone numbers) because the catalogue holds none and an agent should not be nudged into hunting for it. "password" matched, so execute_tool did not raise and never queried MySQL. It returned {"ok": false, "data": null, "error": "blocked by safety rule (sensitive-data): ..."}.

Why it stopped. The harness logged the call and its blocked result, then saw the error begin with "blocked by safety rule". It treats that as a hard stop, not something to retry: it set stop_reason = "safety_block", recorded the message as the answer, wrote the final stop record (1 step, 1 tool

call) and returned without asking the model for another turn. So the run ended at step 1 of 6. It was not `max_steps`, and it was not a normal completion, because the model never produced a final answer.

The rule lives in `execute_tool`, not in the prompt, so it holds even when the model is willing to comply, as here. In S4, a natural request, the model refused by itself and the rule was never needed.

AI Use

1. What did you use an AI assistant for, and what did you do yourself?

I told Claude Code to extend the existing HW4 project (not start a new one), to leave everything unpushed until I had checked it, and to produce a report in the same format as my HW4 report with realistic VS Code-style code images and real output. The domain (Course Catalogue), the FastAPI + MySQL service and the local Ollama model (qwen3:8b) all carry over from earlier homeworks.

Claude Code wrote and ran the implementation: the instructors table and migration, the Pydantic-validated catalog API, the Redux Toolkit client, both MCP servers, the retry/timeout module, the fault-injection experiment, `execute_tool`, the safety rule, the agent loop, the offline tests, `verify_hw05.py`, and the report generator. It drove the real MCP Inspector and a real browser for the screenshots. I am responsible for reviewing the code, the numbers and the report before pushing and submitting.

2. One AI-produced output that was wrong/unsuitable, or one thing you independently verified

The plan assumed the MCP Python SDK's FastMCP class, as the assignment describes. A plain `pip install "mcp[cli]"` now installs MCP 2.x, where FastMCP was renamed, so the first import of `mcp.server.fastmcp` failed and the servers as planned could not run.

3. How you detected the problem or verified the result

The import error was caught immediately by running the import before writing the servers; the SDK's own message named the rename and pointed to pinning `mcp<2`. I also checked that nothing else relied on 2.x APIs by running both servers over real STDIO and through the actual MCP Inspector (mcp dev), not just importing them.

4. What you changed and why it works now

The environment pins `mcp[cli]<2` (v1.30.0), which matches the FastMCP API the assignment is written around. Both servers now start, list the correct tools (4 and 3), and answer valid and invalid calls over STDIO; `verify_hw05.py` repeats that check. A second thing worth recording: my first fault-injection run showed a 139 ms p99 at 0% failures, which was not a retry effect but the first database connection warming up. I added one uncounted warm-up query and re-ran, and the 0% p99 dropped to 11 ms.